



Context

- Requests per second:
 - # of Concurrent requests
 - # of servers
- Response time
 - Autoscaling
 - # of servers



Benchmark

- Client on US West N. California
- System on US West Oregon
 - Connection latency
 - Low overhead



Autoscaling

